

Part 1: Van Gogh Classification

1. Dataset Description

We worked with a **Post-Impressionism subset** of the WikiArt dataset. Each image is labeled by the artist and stored alongside metadata in a CSV file. The classification target was binary:

- 1 if painted by **Vincent van Gogh**
- 0 otherwise

Class Distribution (Train):

Van Gogh: 478

Not Van Gogh: 1522

Class Distribution (Test):

Van Gogh: 197

Not Van Gogh: 1046

This class imbalance highlights the importance of precision and recall-focused metrics like F1 Score and AUC-ROC.

2. Models and Architectures

We fine-tuned **two pretrained models**:

Model	Publication	Original Dataset	Key Features
AlexNet	2012	ImageNet	Simpler, fewer parameters, 5 conv + 3 FC
VGG19	2014	ImageNet	Deeper, 19 layers, very strong at feature extraction

Both models had their classifier head replaced with a single-node linear layer (`nn.Linear(num_features, 1)`) for binary classification.

3. Training Approach

- **Data Augmentation:**
Applied RandomCrop, Horizontal/Vertical Flip, Rotation, ColorJitter.
- **Optimization:**
Optimizer: Adam
Scheduler: ReduceLROnPlateau
Loss: BCEWithLogitsLoss
- **Cross-Validation:**
5-Fold Stratified K-Fold used to improve generalization.

- **Hyperparameter Tuning:**

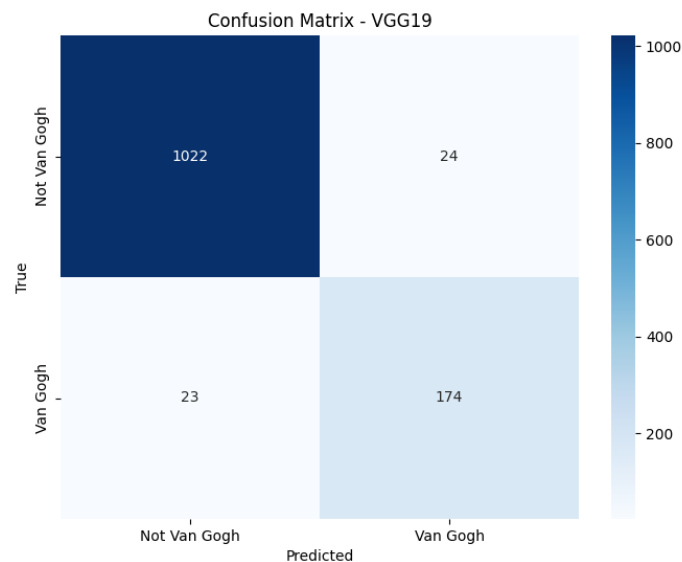
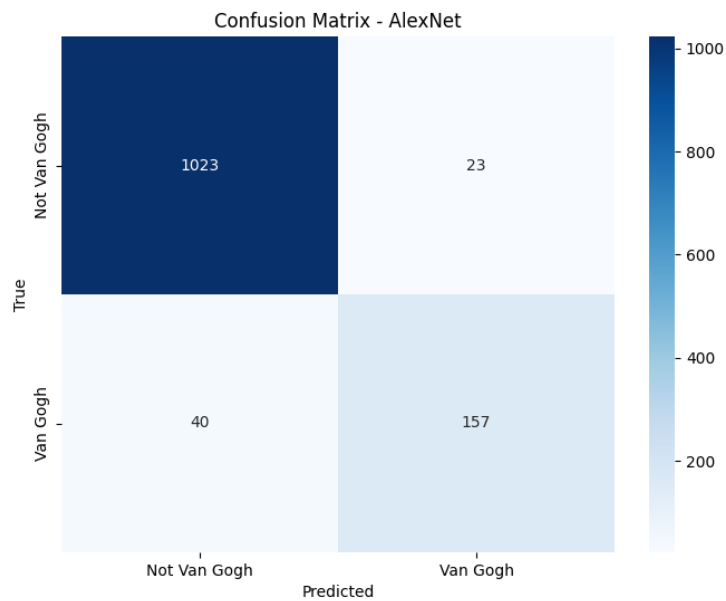
Used **Optuna** to tune:

- Batch size
- Learning rate
- Weight decay

- **Experiment Tracking:**

All key metrics (Accuracy, Loss, F1, AUC) were plotted and logged.

4. Evaluation Metrics and Comparison



- VGG19 significantly reduced false negatives, which is important for identifying Van Gogh paintings.

Metric	AlexNet	VGG19
Best Accuracy	~94.5%	~97.1%
Best F1 Score	~0.83	~0.91
Best AUC-ROC	0.9704	0.9833

VGG19 consistently outperformed AlexNet across all metrics.

5. Metrics Analysis

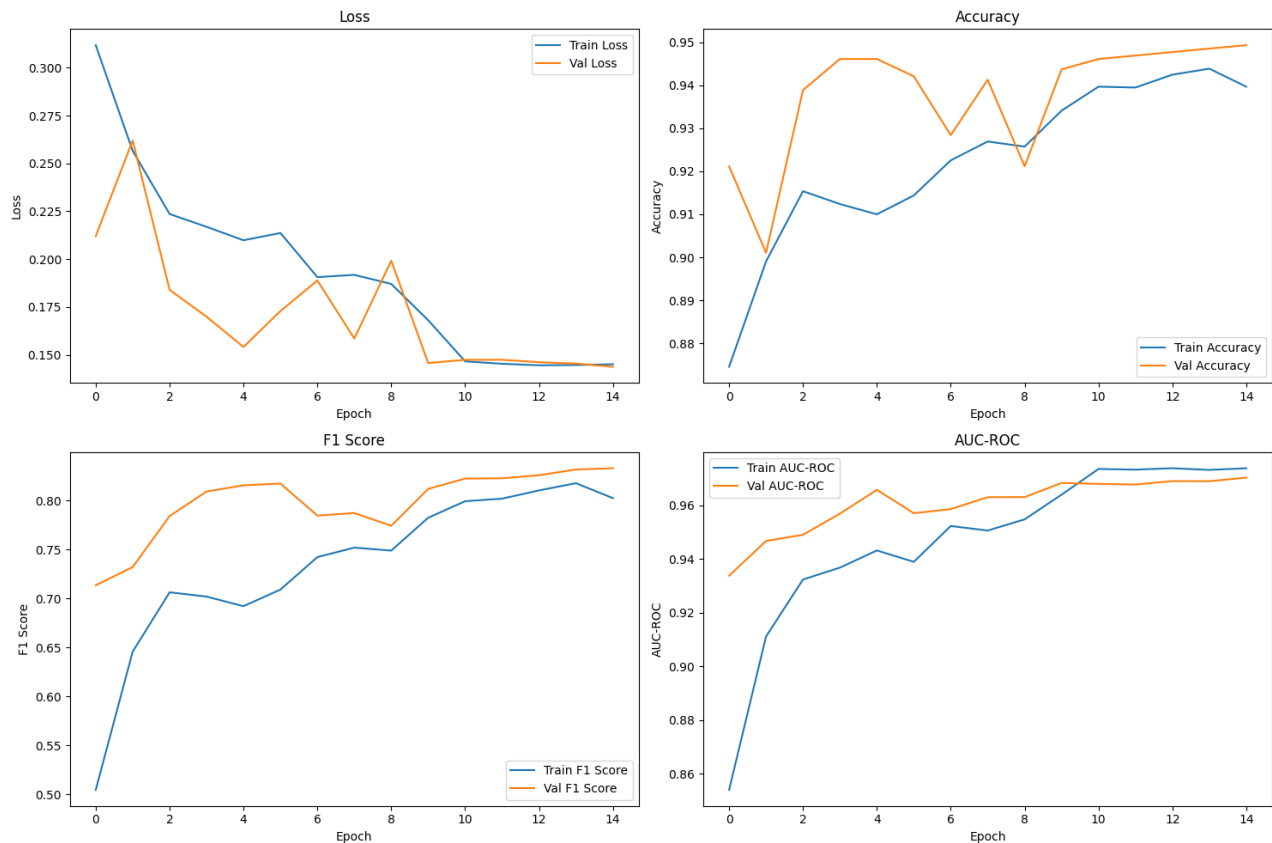
AlexNet

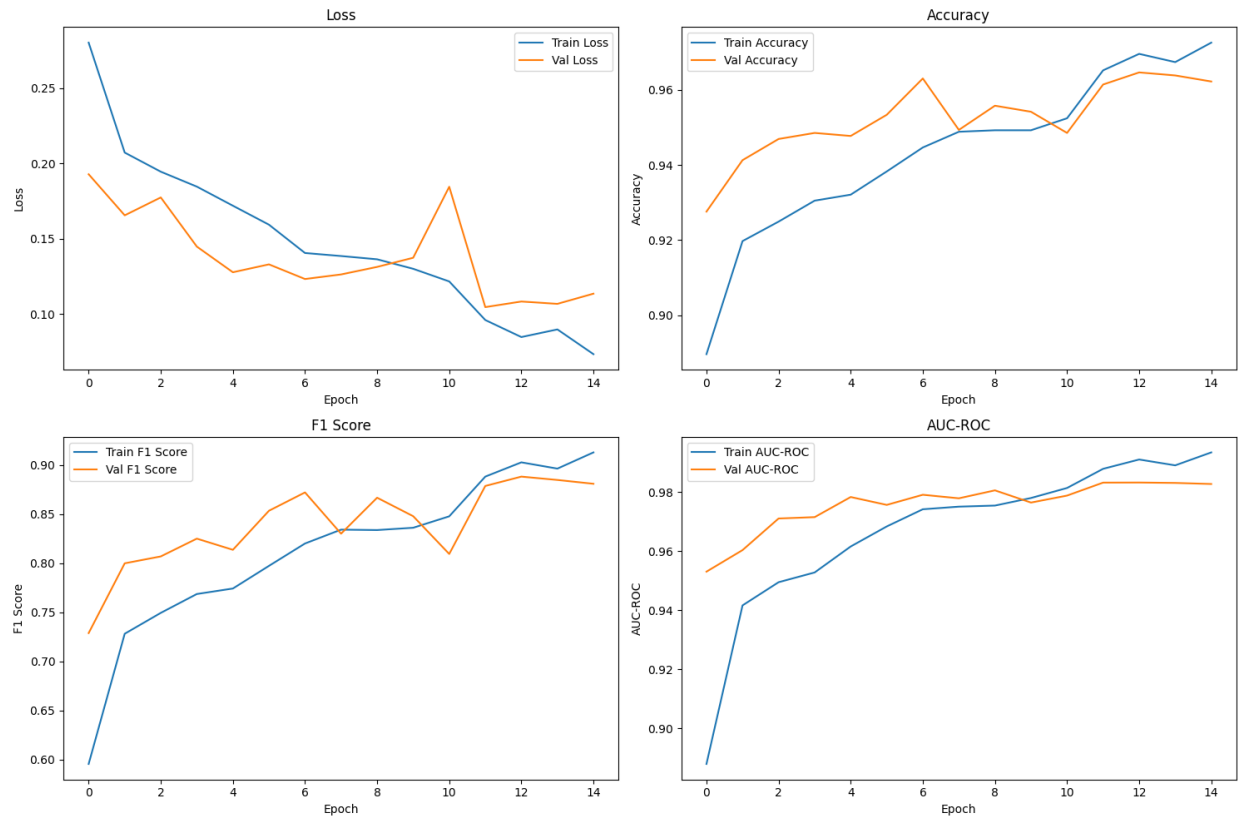
- **False Positives:** 23
- **False Negatives:** 40

VGG19

- **False Positives:** 24
- **False Negatives:** 23

Observation: VGG19 is more balanced, while AlexNet misses more true Van Gogh paintings (higher FN).





6. Model Predictions

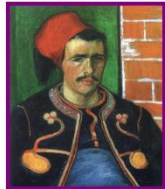
Artist: vincent van gogh
Predicted: Not Van Gogh (12.26%)
✗ Incorrect



Artist: pierre bonnard
Predicted: Not Van Gogh (41.48%)
✓ Correct



Artist: vincent van gogh
Predicted: Not Van Gogh (12.31%)
✗ Incorrect



Artist: vincent van gogh
Predicted: Not Van Gogh (16.23%)
✗ Incorrect



Artist: georges seurat
Predicted: Not Van Gogh (40.99%)
✓ Correct



Artist: theo van rysselberghe
Predicted: Van Gogh (73.40%)
✗ Incorrect



Artist: vincent van gogh
Predicted: Van Gogh (98.27%)
✓ Correct



Artist: vincent van gogh
Predicted: Van Gogh (97.21%)
✓ Correct



Artist: suzanne valadon
Predicted: Not Van Gogh (0.00%)
✓ Correct



Artist: georges seurat
Predicted: Not Van Gogh (0.26%)
✓ Correct



Artist: vincent van gogh
Predicted: Van Gogh (95.89%)
✓ Correct



Artist: vincent van gogh
Predicted: Van Gogh (70.75%)
✓ Correct



Artist: paul cezanne
Predicted: Not Van Gogh (0.00%)
✓ Correct



Artist: paul cezanne
Predicted: Not Van Gogh (0.02%)
✓ Correct



Part 2: Style Transfer

1. Style Transfer Function

- Built a **generic PyTorch function** compatible with both **VGG19** and **AlexNet**.
- Combines:
 - **Content Loss** from high-level layers (e.g. conv_5)
 - **Style Loss** using **Gram matrices** from multiple low- and mid-level layers.
- Style weights: [1.0, 0.8, 0.6, 0.4, 0.2]
- Style intensity (style_weight) = 1e6
- Content intensity (content_weight) = 1

2. Input Details

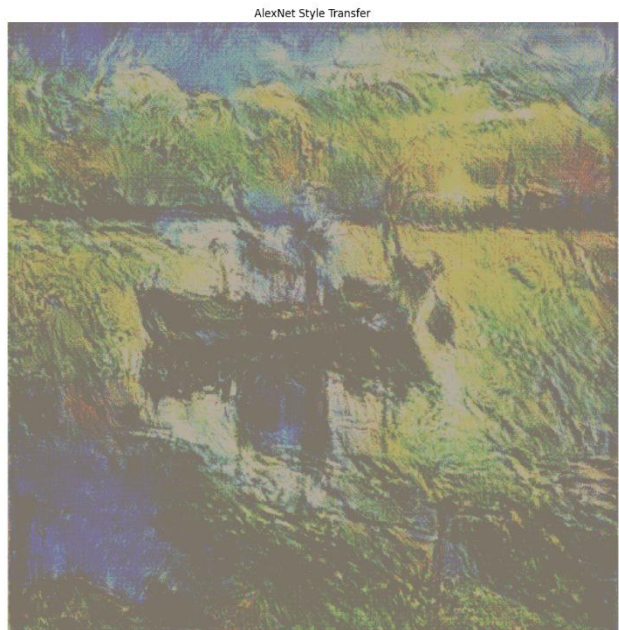
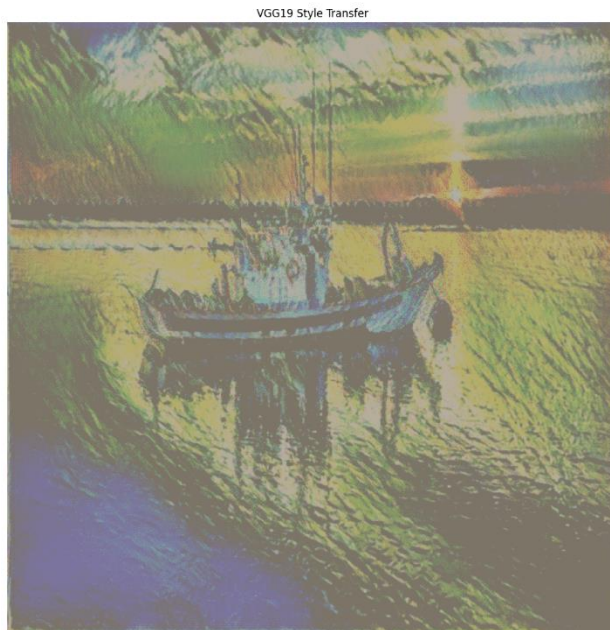
- **Content Images:** 1 personal photo
- **Style Images:** 5 Van Gogh paintings
- **Epochs:** 300 (fixed for all)
- **Output:** Stylized images from both VGG19 and AlexNet

3. Generated Paintings

Base Painting:



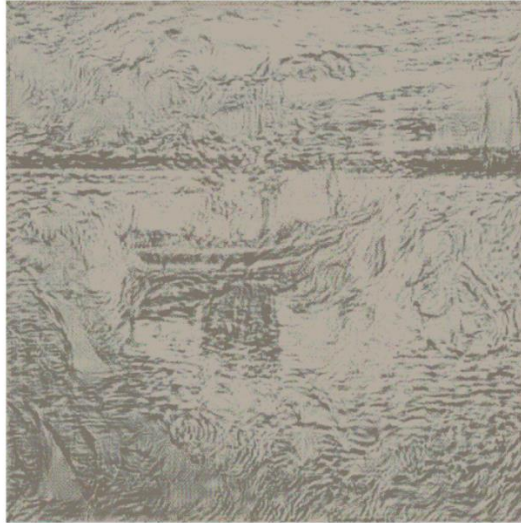
Style Transferred Paintings:



VGG19 Style Transfer



AlexNet Style Transfer



VGG19 Style Transfer



AlexNet Style Transfer



VGG19 Style Transfer



AlexNet Style Transfer



VGG19 Style Transfer



AlexNet Style Transfer



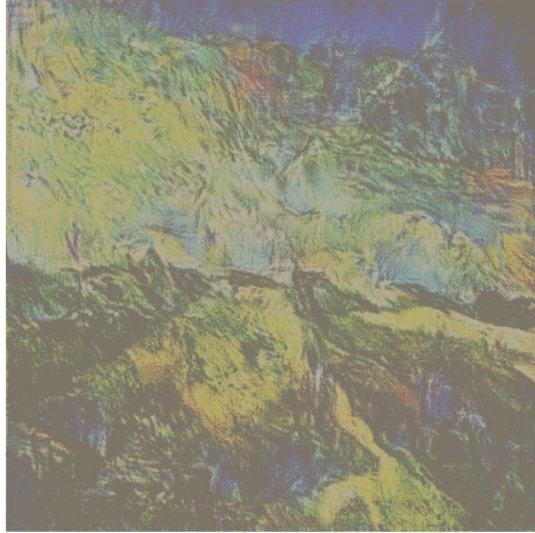
Another Example:



VGG19 Style Transfer



AlexNet Style Transfer



VGG19 Style Transfer



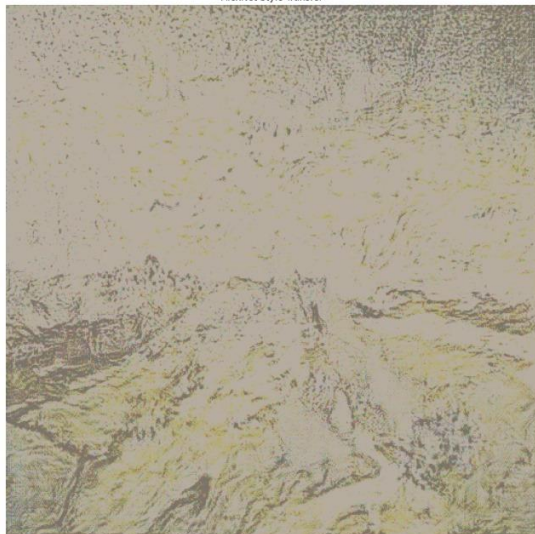
AlexNet Style Transfer

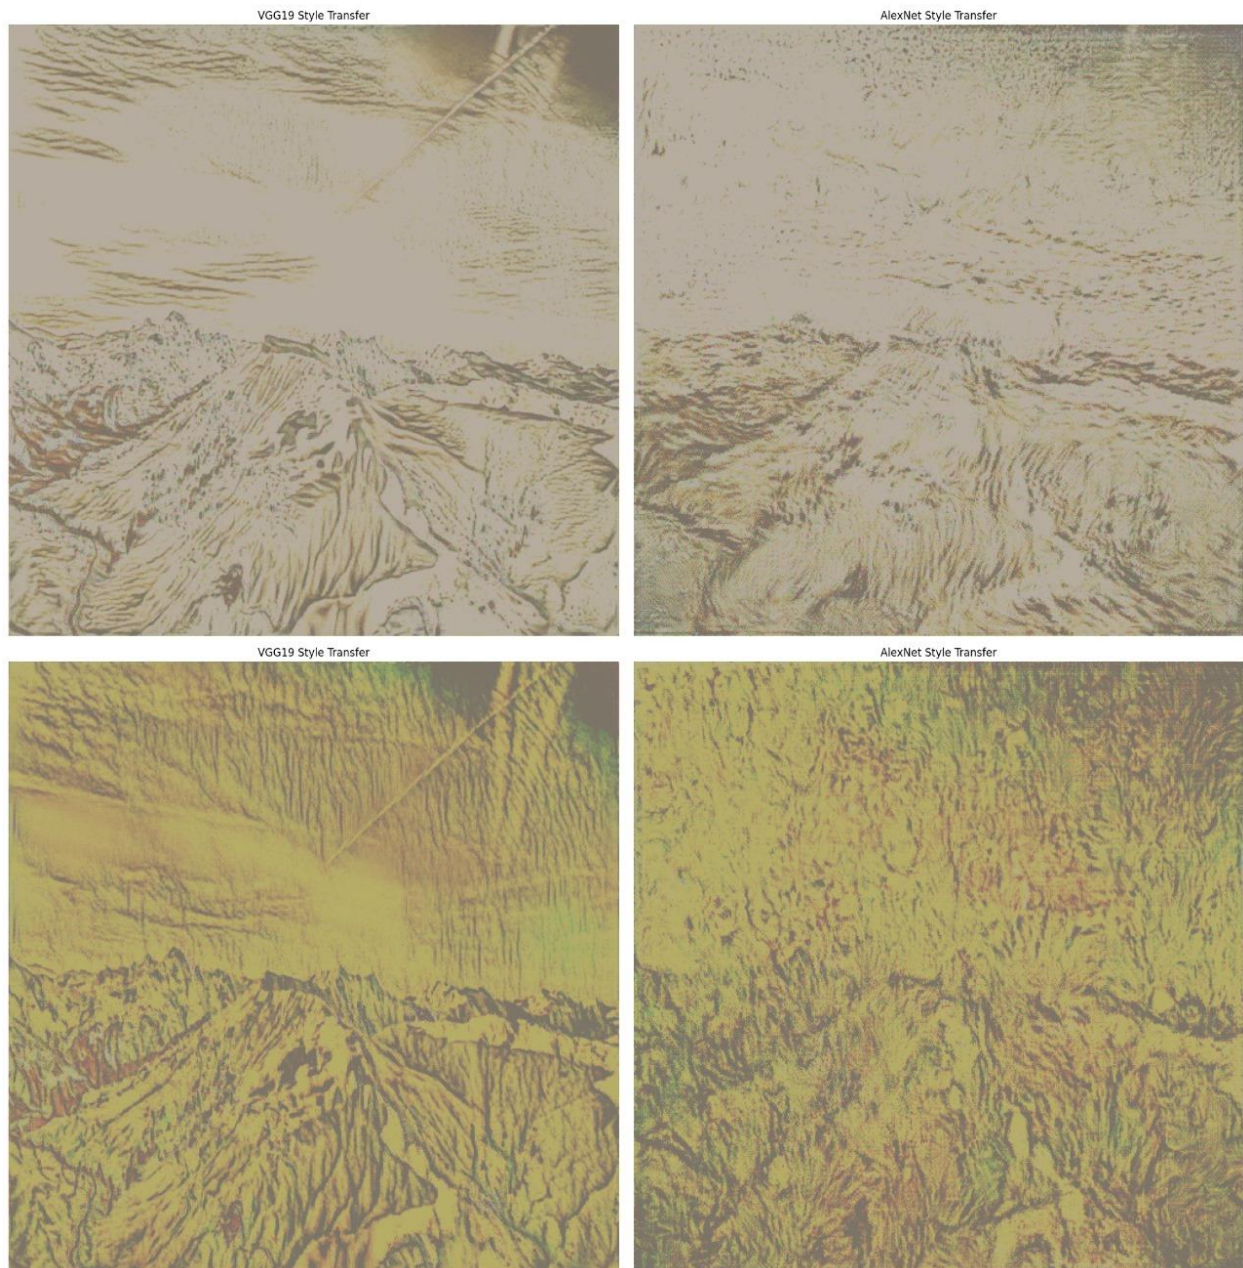


VGG19 Style Transfer



AlexNet Style Transfer





4. Classification-Based Evaluation

Classification Results (% of images predicted as Van Gogh):

Generated By	Classified by VGG19	Classified by AlexNet
VGG19	60%	80%
AlexNet	80%	80%

5. Interpretation

- **VGG19-generated outputs** were more artistically stylized but sometimes less recognizable to the classifier as Van Gogh.
 - **AlexNet-generated outputs** seemed more “conservative”, retaining more recognizable texture/patterns for the classifiers.
 - **Classifier agreement** was higher on AlexNet-style results.
-

5. Layer Choices

Model	Content Layer	Style Layers
VGG19	conv_5	conv_1 to conv_5
AlexNet	conv_5	conv_1 to conv_5

Justification:

- Shallow layers capture texture (style)
 - Deeper layers preserve object layout (content)
-

6. Validation Through Classification

The classification models partially aligned with human perception:

- When both classifiers agree on Van Gogh style, the result is clearly artistic.
- Misalignments reveal that classifier "style recognition" doesn't always align with human aesthetic judgment.

Code Explanation:

Part 01:

1. Imports and Setup

- **Libraries:** Uses PyTorch, torchvision, pandas, scikit-learn, Optuna (for hyperparameter tuning), and matplotlib/seaborn for visualization.
- **Paths:** Defines dataset paths (PAINTINGS_PATH for images, CSV_PATH for labels).
- **Device:** Uses GPU if available.
- **Reproducibility:** Sets random seeds for PyTorch, NumPy, and CUDA.

2. Custom Dataset (WikiArtDataset)

- **Purpose:** Loads images and labels from a CSV file, filters by subset (train/test), and checks image validity.
- **Key Features:**
 - **Label Creation:** Marks images as Van Gogh (1) if "van gogh" is in the artist's name (case-insensitive).
 - **Validation:** Filters out invalid image paths to avoid runtime errors.
 - **Transforms:** Applies data augmentation (training) or standard preprocessing (validation).

3. Data Augmentation/Preprocessing (get_transforms)

- **Training Transforms:** Random resizing, flipping, rotation, and color jitter to prevent overfitting.
- **Validation Transforms:** Fixed resizing and center cropping for consistent evaluation.
- **Normalization:** Uses ImageNet mean/std for compatibility with pre-trained models.

4. Model Architecture (FineTunedModel)

- **Base Models:** Supports VGG19 or AlexNet from torchvision.models.
- **Transfer Learning:** Replaces the final classification layer with a single-output neuron for binary classification (Van Gogh vs. not).
- **Output:** Raw logits (passed through sigmoid during evaluation for probabilities).

5. Training Loop (train_model)

- **Training Phase:** Computes loss, updates weights via backpropagation, and tracks metrics (loss, accuracy, F1, AUC).
- **Validation Phase:** Evaluates model performance on a held-out set.
- **Early Stopping:** Uses ReduceLROnPlateau to adjust learning rates if validation loss stagnates.
- **Checkpointing:** Saves the model with the best validation AUC.

6. K-Fold Cross-Validation (perform_kfold)

- **Stratified Splits:** Maintains class balance across folds using StratifiedKFold.

- **Training per Fold:** Trains the model on each fold and records validation AUC.
- **Result Aggregation:** Reports average AUC across all folds.

7. Hyperparameter Tuning (objective)

- **Optuna Integration:** Optimizes hyperparameters like:
 - Batch size (8, 16, 32)
 - Learning rate (log-scale between $1e-5$ and $1e-3$)
 - Weight decay (L2 regularization).
- **Short Training:** Runs 5 epochs per trial to balance speed and performance.

8. Visualization Tools

- **Metric Plots:** Generates loss, accuracy, F1, and AUC curves (`plot_metrics`).
- **Confusion Matrix:** Visualizes true vs. predicted labels for the test set (`plot_confusion_matrix`).

9. Workflow (main)

1. **Data Loading:** Validates image paths and checks class distribution.
2. **Hyperparameter Tuning:** Runs Optuna trials for VGG19 and AlexNet.
3. **Cross-Validation:** Evaluates model stability across 5 folds.
4. **Final Training:** Trains models with optimal hyperparameters for 15 epochs.
5. **Evaluation:** Plots metrics and confusion matrices for the test set.

Part 02:

Key Components

1. **Image Loading/Preprocessing:**
 - `load_image`: Resizes and normalizes images using ImageNet statistics (for compatibility with pre-trained models).
 - `im_convert`: Converts tensors back to numpy arrays for visualization.
2. **Loss Functions:**
 - **Content Loss:** Measures the MSE between intermediate feature maps of the content image and the generated image.
 - **Style Loss:** Uses **Gram matrices** to compare the correlation of feature maps between the style image and generated image. This captures texture/style rather than content.

3. **Model Setup:**

- `get_model_and_losses`: Builds a modified VGG19/AlexNet model with hooks to compute content/style losses at specified layers.
- **Content Layers**: Typically deeper layers (e.g., `conv_5`) to preserve content.
- **Style Layers**: Shallow and deep layers (e.g., `conv_1` to `conv_5`) to capture style at multiple scales.

4. **Optimization:**

- Uses **LBFGS optimizer** to iteratively adjust the generated image to minimize the weighted sum of content and style losses.
- Parameters like `style_weight` and `content_weight` control the trade-off between preserving content vs. style.

5. **Style Transfer Execution:**

- `neural_style_transfer`: Runs the optimization loop to generate the stylized image.
- `compare_style_transfer`: Compares results from VGG19 and AlexNet side-by-side.

6. **Batch Processing:**

- `batch_style_transfer`: Applies style transfer to multiple content/style image pairs and saves results.

7. **Evaluation:**

- `evaluate_style_transfer`: Uses the **classification models** from the first script to check if the generated images are classified as "Van Gogh style" (measuring style transfer effectiveness).

Workflow

1. **Input Preparation:**

- Load content (e.g., a photo) and style images (e.g., Van Gogh's painting).
- Preprocess images into tensors.

2. **Model Configuration:**

- Choose a backbone (VGG19/AlexNet).
- Define which layers to use for content/style loss and their weights.

3. **Optimization Loop:**

- Forward pass: Compute content/style losses.

- Backward pass: Update the generated image to minimize losses.
- Repeat for num_epochs iterations.

4. **Result Visualization:**

- Display/save the stylized image and compare VGG19/AlexNet outputs.